



Shorter Messages and Faster Post-Quantum Encryption with Round5 on Cortex M

Markku-Juhani O. Saarinen <mjos@pqshield.com>

S. Bhattacharya¹ O. Garcia-Morchon¹ R. Rietman¹ L. Tolhuizen¹ Z. Zhang²

(¹)Philips Research, NL (²)Algorand, USA (Prev. OnBoard Security)

CARDIS 2018

13 November 2018 – Montpellier, France



UNIVERSITY OF
OXFORD



OXFORD SCIENCES
INNOVATION

Innovate UK
Technology Strategy Board

Round2 + Hila5 = Round5

Dustin Moody (NIST) on April 30, 2018:

“NIST would like to encourage any submissions which are quite similar to consider merging. It would be helpful if any such merger is announced (to NIST) before November 30th. [...] While the selection of candidates for the second round will primarily be based on the original submissions, NIST may consider a merged submission more attractive than either of the original schemes if it provides improvements in security, efficiency, or compactness and generality of presentation.”

- ▶ We took the NIST advice and decided to merge my personal round 1 candidate **Hila5** with **Round2**, primarily from Philips. The resulting **Round5** algorithm has better security analysis and is much more efficient than either one of its parents.
- ▶ Round5 resembles the ring variants of Round2 more than Hila5. Hila5 was based on Ring-LWE in anticyclic $x^n + 1$ “NewHope” ring ($n = 1024$, $q = 12289$), but with a novel reconciliation method - and importantly **XEf error correction codes**.

General and Ideal Lattices


Lattice basis:

$$\mathcal{L}(b_1, \dots, b_n) = \sum_{i=1}^n x_i b_i \mid x_i \in \mathbb{Z}$$

If we replace b_2 with the shorter green vector, same lattice is generated. This is a **reduced basis**.

Shortest Vector Problem (**SVP**) is the task of finding the shortest vector in a lattice. Related hard lattice problems include: Bounded Distance Decoding (**BDD**), Short Integer Solution (**SIS**), Learning With Errors (**LWE**), and Learning With Rounding (**LWR**).

In order to go from n^2 variables to n variables, one may use **ideal lattices**, which are defined in **rings**.



Lattice \mathbb{Z}^2 with basis (b_1, b_2) .

(G)LWR vs (R/M)LWE

Hila5, Kyber, NewHope, Frodo, etc. are based on versions of the *Learning With Errors* (**LWE**) problem, which requires addition of “noise” sampled from a bell-shaped (e.g. Gaussian or Binomial) distribution to create hard instances.

(G)LWR: Noise from Rounding (no RNG)

Round5, Round2, Saber instead use versions of *Learning With Rounding* (**LWR**). No sampling is required in LWR; “noise” comes directly from rounding in the form of a lossy compression function we may call **Round**. It maps $x \in \mathbb{Z}_a$ to \mathbb{Z}_b using h :

$$\text{Round}_{a \rightarrow b}(x, h) = \left\lfloor \frac{b}{a} \cdot x + h \right\rfloor \bmod b. \quad (1)$$

This is equivalent to rounding of $b/a \cdot x$ to closest integer when rounding constant $h = 1/2$. Each coefficient is operated on separately when **Round** or modular reduction (“**mod**”) is applied to polynomials, vectors, or matrices.

Hila5: Ring-LWE (Learning With Errors – in a Ring)

Hila5 used the Ring Learning With Errors (RLWE) problem. Let \mathcal{R} be a ring with elements $\mathbf{v} \in \mathbb{Z}_q^n$. As usual, $n = 2^m$ power of two and $n \mid q - 1$ for **NTT** to work.

Definition (RLWE – Informal)

With all distributions and computations in ring \mathcal{R} , let \mathbf{s}, \mathbf{e} be elements randomly chosen from some non-uniform distribution χ , and \mathbf{g} be a uniformly random public value. Determining \mathbf{s} from $(\mathbf{g}, \mathbf{g} * \mathbf{s} + \mathbf{e})$ in ring \mathcal{R} is the (Normal Form Search) Ring Learning With Errors ($\text{RLWE}_{\mathcal{R}, \chi}$) problem.

The hardness of the problem is a function of n , q , and χ . Typically χ is chosen so that each coefficient is a Gaussian or from some other “Bell-Shaped” distribution that is relatively tightly concentrated around zero. Hila5 used a binomial “bitcount” χ .

Round2, Round5: General Learning With Rounding

Definition (General LWR (GLWR) – Informal)

Let d, n, p, q be positive integers such that $q \geq p \geq 2$, and $n \in \{1, d\}$. Let $\mathcal{R}_{n,q}$ be a polynomial ring, and let D_s be a probability distribution on $\mathcal{R}_n^{d/n}$.

- ▶ The search version of the GLWR problem $\text{sGLWR}_{d,n,m,q,p}(D_s)$ is as follows: given m samples of the form $(\mathbf{a}_i, b_i = \text{Round}_{q \rightarrow p}(\mathbf{a}_i^T \mathbf{s} \bmod q, 1/2))$ with $\mathbf{a}_i \in \mathcal{R}_{n,q}^{d/n}$ and a fixed $\mathbf{s} \leftarrow D_s$, recover \mathbf{s} .
- ▶ The decision GLWR problem $\text{dGLWR}_{d,n,m,q,p}(D_s)$ is to distinguish between the uniform distribution on $\mathcal{R}_{n,q}^{d/n} \times \mathcal{R}_{n,p}$ and the distribution $(\mathbf{a}_i, b_i = \text{Round}_{q \rightarrow p}(\mathbf{a}_i^T \mathbf{s}_i \bmod q, 1/2))$ with $\mathbf{a}_i \leftarrow \mathcal{R}_{n,q}^{d/n}$ and a fixed $\mathbf{s} \leftarrow D_s$.

Here $n = d$ implies the ring variant that we are using in Round5. Typically p, q are powers of two (no modular reduction) and there is more freedom in choosing n .

Round5 rough high-level description: Based on “Noisy ElGamal”

- ▶ Round5 comes in **both** ring and non-ring variants – in this work we will focus on faster and smaller ring ($n = d$) variants on low-end microcontrollers (Cortex M).
- ▶ **XEf** is a forward error correction code that can be easily implemented in constant time fashion (this comes from Hila5).
- ▶ **Two rings** are used: $x^{n+1} - 1$, with $n + 1$ prime (cyclic, a bit like HILA5), and its subring $\Phi_{n+1} = (x^{n+1} - 1)/(x - 1) = x^n + \dots + x + 1$ (resembling Round2).
- ▶ For small-norm secrets we use “balanced” sparse ternary polynomials $D \subset \{-1, 0, 1\}^n$, with $h/2$ of both $+1$ and -1 coefficients ($n - h$ set to 0).
- ▶ The CPA versions **R5ND_nKEM** are intended for key establishment while CCA versions **R5ND_nPKE** are for public key encryption. $n \in \{1, 3, 5\}$ is the NIST security level (corresponding to security of AES with $\{128, 192, 256\}$ bit key).
- ▶ The CCA KEM can be used as Public Key Encryption Algorithm with a DEM: SHAKE-256 to expands the shared secret to AES-GCM key and IV for data.

Round5 Ring Variants: Simple Key Generation

KeyGenCPA(σ, γ): Key generation for CPA case.

Require: Random seeds σ, γ .

1: $\mathbf{a} \xleftarrow{\$ \sigma} \mathbb{Z}_q^n$

Uniform polynomial, seed σ .

2: $\mathbf{s} \xleftarrow{\$ \gamma} D$

Sparse ternary polynomial, seed γ .

3: $\mathbf{b} \leftarrow \text{Round}_{q \rightarrow p}(\mathbf{a} * \mathbf{s} \bmod \Phi_{n+1}, 1/2)$

Compress product to range $0 \leq b_i < p$.

4: $\mathbf{sk} = \mathbf{s}$

\mathbf{s} : Random seed γ is sufficient.

5: $\mathbf{pk} = (\mathbf{a}, \mathbf{b})$

\mathbf{a} : Random seed σ , \mathbf{b} : $n \log_2 p$ bits.

Ensure: Public key $\mathbf{pk} = (\mathbf{a}, \mathbf{b})$ and secret key $\mathbf{sk} = \mathbf{s}$.

Note: This description is simplified; e.g. the rounding constants are not always $1/2$.

Round5 Ring Variants: Encryption uses (XEf) Error Correction

EncryptCPA(m, pk, ρ): Public key encryption (CPA).

Require: Message $m = \{0, 1\}^m$, public key $\mathbf{pk} = (a, b)$, random seed ρ .

1: $r \xleftarrow{\$ \rho} D$

Sparse ternary polynomial, seed ρ

2: $u \leftarrow \text{Round}_{q \rightarrow p}(a * r \bmod \Phi_{n+1}, 1/2)$

Compress product to range $0 \leq u_i < p$.

3: $t \leftarrow \text{Sample}_{\mu}(b * r \bmod x^{n+1} - 1)$

Noisy shared secret, truncate to \mathbb{Z}_p^{μ} .

4: $v \leftarrow \text{Round}_{p \rightarrow t}(t + \frac{p}{2}m, 1/2)$

Add message + XEf $m \in \mathbb{Z}_2^{\mu}$.

5: $\mathbf{ct} = (u, v)$

u : $n \log_2 p$ bits, v : $\mu \log_2 t$ bits.

Ensure: Ciphertext $\mathbf{ct} = (u, v)$.

The CPA scheme is transformed into a chosen-ciphertext (IND-CCA2) secure one (**R5ND_xPKEb**) using the usual Fujisaki-Okamoto Transform [FuOk99, HoHoKi17].

Decryption

DecryptCPA(ct, sk): Decryption (CPA)


Require: Ciphertext $\mathbf{ct} = (\mathbf{u}, \mathbf{v})$, secret key $\mathbf{sk} = \mathbf{s}$.

- 1: $\mathbf{t}' \leftarrow \text{Sample}_\mu(\mathbf{u} * \mathbf{s} \bmod x^{n+1} - 1)$ *Noisy shared secret, truncate to \mathbb{Z}_p^μ .*
- 2: $\mathbf{m} \leftarrow \text{Round}_{p \rightarrow 2}(\frac{p}{t} \mathbf{v} - \mathbf{t}', 1/2)$ *Remove noise, apply XEf correction.*

Ensure: Plaintext $\mathbf{pt} = \mathbf{m}$.

- ▶ To see why the algorithm works, note that the shared secrets **EncryptCPA()** and **DecryptCPA()** satisfy approximately $\mathbf{t} \approx \mathbf{t}' \approx \mathbf{a} * \mathbf{s} * \mathbf{r}$.
- ▶ The Φ_{n+1} ring a subring of the $x^{n+1} - 1 = (x - 1) * \Phi_{n+1}$ ring. Since \mathbf{s} and \mathbf{r} are “balanced”, their coefficients sum to zero and they are divisible by $(x - 1)$.
- ▶ High bits of \mathbf{t} are used as a “one time pad” to transport the message payload.

Error Correction Code XEf



← Hey kids! Pay attention in coding theory classes. It's useful.

I chose a simple forward error correction code, XE5, for HILA5. Error correction helps to significantly shrink message sizes.

Additional Requirement: Fast, constant-time implementable.

XE5 has been generalized as **XEf** for Round5, where $f \in \{1, 2, 3, 4, 5\}$ is the number of bit flips it can always correct.

- ▶ **XEf** is a simple parity code. Both encoding and decoding are parallelizable and easy to implement efficiently on 8 to 64-bit architectures and in hardware.
- ▶ **Constant-time.** No table lookups or branches. An “error oracle” would lead to attacks under CCA security model. The KEMs also never fail (QRom proofs).

I first described similar constant-time error correction techniques (for TRUNC8) in: <https://eprint.iacr.org/2016/1058> (Original uploaded November 15, 2016)

Round5 Ring Variants for Key Establishment

	Parameters	CPA NIST1	CPA NIST3	CPA NIST5
Round5.KEM	d, n, h	490, 490, 162	756, 756, 242	940, 940, 414
	q, p, t	$2^{10}, 2^7, 2^4$	$2^{12}, 2^8, 2^2$	$2^{12}, 2^8, 2^3$
	B, \bar{n}, \bar{m}, f	1, 1, 1, 3	1, 1, 1, 3	1, 1, 1, 3
	μ	128 + 91	192 + 103	256 + 121
	Public key	445 B	780 B	972 B
	Ciphertext	539 B	830 B	1082 B
	PQ Security	2^{118}	2^{176}	2^{232}
	Classical	2^{128}	2^{193}	2^{256}
	Failure rate	2^{-78}	2^{-78}	2^{-95}
	Version ID	R5ND_1KEMb	R5ND_3KEMb	R5ND_5KEMb

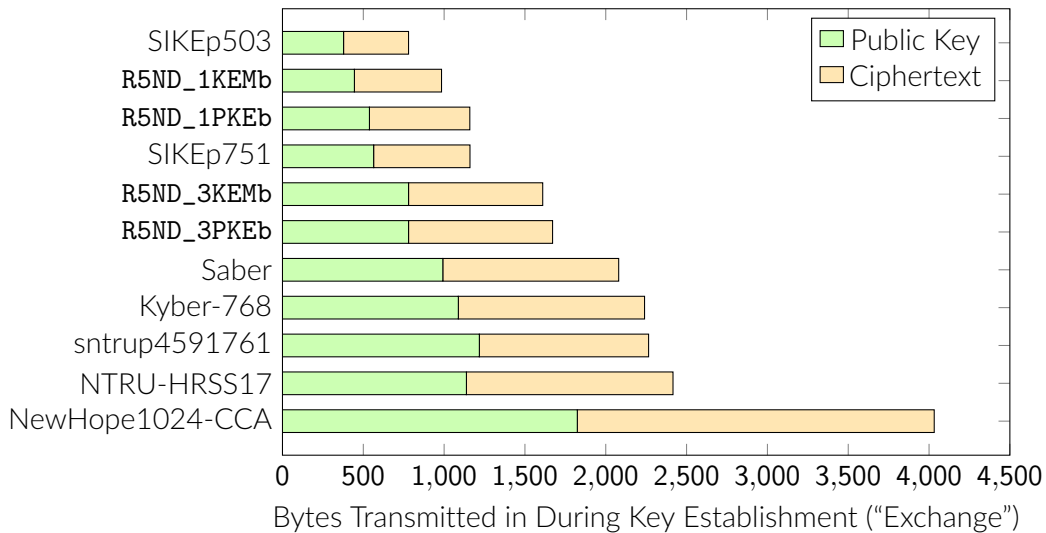
In TLS, SSH, ... Key Establishment, ephemeral keys are used, so CPA is fine. Also a higher failure probability is acceptable since the process can be simply repeated.

Round5 Ring Variants for Public Key Encryption

	Parameters	CCA NIST1	CCA NIST3	CCA NIST5
Round5.PKE	d, n, h	522, 522, 208	756, 756, 242	940, 940, 406
	q, p, t	$2^{13}, 2^8, 2^3$	$2^{12}, 2^8, 2^3$	$2^{12}, 2^8, 2^4$
	B, \bar{n}, \bar{m}, f	1, 1, 1, 3	1, 1, 1, 3	1, 1, 1, 3
	μ	128 + 91	192 + 103	256 + 121
	Public key	538 B	780 B	972 B
	Ciphertext	621 B	891 B	1161 B
	PQ Security	2^{117}	2^{176}	2^{232}
	Classical	2^{128}	2^{193}	2^{256}
	Failure rate	2^{-202}	2^{-171}	2^{-131}
	Version ID	R5ND_1PKEb	R5ND_3PKEb	R5ND_5PKEb

Public Key Encryption variants have IND-CCA2 security (long-term keys) and have a negligible failure probability. You can use these e.g. for e-mail or cloud backups.

Parameters are Optimized for Bandwidth



Polynomial Multiplication with Sparse Ternary Secrets

- ▶ Polynomial multiplication (and reduction) with our sparse ternary secrets of weight h degree n requires $\Theta(nh)$ word additions / subtractions.
- ▶ Even though h is a lot smaller than n , and no word multiplications are needed, this is essentially an $O(n^2)$ algorithm.
- ▶ Since p, q, t are powers of two, no modular reduction is needed. So the complexity of this simple multiplication algorithm is $\Theta(nh \log_2 q)$ bit operations.
- ▶ Also **no word multiplications** are required; this is useful on extremely low-end CPUs such as 8-bit AVR that does not have a word \times word multiplier.
- ▶ Parallel byte or word additions / subtractions resemble string operations. They are also directly supported by **SIMD** architectures (AVX2, AVX-512, ARM NEON) which makes them fast on high-end CPUs (and in **hardware**).
- ▶ In other words, the special form allows smaller implementation footprint in both hardware and software. Flexible, fast implementations are **easy to write**.

Polynomial Multiplication for General, Non-Ternary Secrets

This is much, much harder [1]:

“[...] systematically exploring different combinations of Toom-3, Toom-4, and Karatsuba decomposition of multiplication in \mathcal{R}_q , and by carefully hand-optimizing multiplication of low-degree polynomial multiplication at the bottom of the Toom/Karatsuba decomposition.”

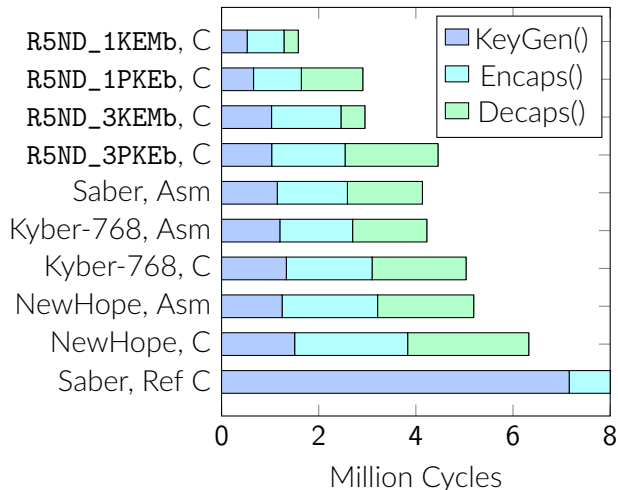
[1] M. J. Kannwischer, J. Rijneveld and P. Schwabe: “*Faster multiplication in $\mathbb{Z}_{2^m}[\mathbf{x}]$ on Cortex-M4 to speed up NIST PQC candidates*”, <https://eprint.iacr.org/2018/1018>, Oct 2018.

- ▶ NTT: $\Theta(n \log n)$ – requires smooth n and $n \mid q - 1$.
- ▶ Toom-3: $\Theta(n^{\frac{\log 5}{\log 3}}) \approx \Theta(n^{1.46497})$
- ▶ Karatsuba: $\Theta(n^{\log_2 3}) \approx \Theta(n^{1.58496})$, etc.


$\Theta(1)$	=	$\Theta(\text{😎})$
$\Theta(\log(n))$	=	$\Theta(\text{😊})$
$\Theta((\log(n))^c)$	=	$\Theta(\text{😌})$
$\Theta(n)$	=	$\Theta(\text{🙂})$
$\Theta(n \log(n))$	=	$\Theta(\text{😊})$
$\Theta(n^{1.5})$	=	$\Theta(\text{😐})$
$\Theta(n^2)$	=	$\Theta(\text{😞})$
$\Theta(n^c)$	=	$\Theta(\text{😡})$
$\Theta(c^n)$	=	$\Theta(\text{😱})$
$\Theta(n!)$	=	$\Theta(\text{😨})$

The asymptotically best algorithm is not necessarily the fastest for given n .

Performance of our C implementation on Cortex M4



"NewHope" refers to NewHope1024CCA. Saber is [KaMeRo+18].



My \$20 Teensy 3.2 dev board has a (\$5) **MK20DX256** @ 96 MHz.
"Hi-spec": 64kB RAM, 256kB Flash.
(A lo-spec Cortex M4 MCU costs \$1.)

Performing Key Establishment at NIST 3 Security Level

Xfer: Public key + Ciphertext. **Time:** KeyGen + Encaps + Decaps on M4 @ 24 MHz.

Code: Size of implementation in bytes. **Fail:** Decryption failure bound.

PQ: Claimed quantum security. **Classic:** Claimed classical security.

Algorithm	Xfer	Time	Code	Fail	PQ	Classic
R5ND_3KEMb (C)	1610	0.123s	4464	2^{-78}	2^{176}	2^{193}
R5ND_3PKEb (C)	1671	0.185s	5232	2^{-171}	2^{176}	2^{193}
Saber (CHES18 Asm)	2080	0.172s	?	2^{-136}	2^{180}	2^{198}
Kyber-768 (Asm)	2240	0.210s	7016	2^{-142}	2^{161}	2^{178}
sntrup4591761 (C)	2265	8.718s	71024	0	?	2^{248}
NTRU-HRSS17 (C)	2416	7.814s	11956	0	2^{123}	2^{136}
NewHope1024-CCA	4032	0.264s	12912	2^{-216}	2^{233}	?
SIKEp751 (C)	1160	685.9s	19112	0	2^{124}	2^{186}

Conclusions

- ▶ **Round5** is a merger of two NIST Post-Quantum Competition project KEM & Public Key Encryption algorithms: Hila5 (from me) and Round2 (mainly Philips).
- ▶ **GLWR**. Based on (Generalized) Learning With Rounding – no need for random numbers or random samplers, helps to shrink message sizes.
- ▶ **Large design space** of n (dimension), p , q , t (moduli - just bit masking), h (ternary secret weight), f (error correction), etc. compared to Ring-LWE or Module-LWE candidates. This helped us to **carefully optimize the parameters**.
- ▶ **Bandwidth**. Key sizes and message expansion of Round5 are the smallest among all NIST candidates suitable for embedded use.
- ▶ **Performance**. Even the portable C implementation outperforms most assembler-optimized schemes on embedded and PC. HW profile is also good.

Thank You!